

PROBE-R3 — Findings

E2 v27.7 + E1 v2.7, two runs. Thirteen records against two job descriptions. Winnow pipeline, last probe before deploy.

Run date 2026-09-05 · reported 2026-09-06 · for the master thread

What this is

The short form of PROBE-R3-REPORT (145 pp). Everything here is in that document; nothing here is in this one alone. Where a claim below is a quotation, it is complete. Where it is a count, the full working is in the section named.

Nothing was scored by the session that ran this. The brief's predictions are printed beside what happened; findings are recorded as findings; none was fixed.

Install held. All four hashes match. Both assets were **added** — stage1-bootstrap-v2.7.txt and stage3-bootstrap-v27.7.txt — hash-verified before being manifested, two new manifest entries, no existing entry edited, and both superseded assets still present and still hashing to their manifested values. config.json repointed; no extraction key. intake.py 03 and run.py 05 unchanged. load_verified_assets() silent before and after. No HASH MISMATCH.

All six invocations completed straight through. 108 provider calls, zero failures, zero retries, every one finish_reason: stop on the first attempt.

Two things could not be done

Kittiya is not in this workspace. No résumé, no run directory, no numbered text, under that spelling or any other. The only occurrences of the name anywhere are inside probe-r3.zip's own two markdown files; FRONTEND-FIX-01-honorific-strip.md places her at R4 Ms.Kittiya Anantapatthamanond under **v27.5**, a probe whose inputs are not in winnowpipelineCURRENT20260906.zip. Every candidate present in the workspace was enumerated before concluding this. PROBE-R3A-S2 and PROBE-R3B-S2 therefore carry **Runyarin only** — thirteen records ran, not fourteen. No substitute was chosen and nothing was invented. R3-P1's Kittiya row is unanswerable and is reported as such.

FRONTEND-FIX-01-honorific-strip.md was not acted on. render.py does not exist anywhere in this workspace. The file was read only far enough to confirm it is a frontend spec and to identify Kittiya, and is carried untouched in the evidence bundle for whoever holds render.py.

The headline

- **R3-P1 holds completely.** All twelve named-tool lines are admitted, in both runs. All three generic class-word inventory lines are still NULLED, in both runs. E1 v2.7 does what the ruling asked and nothing more.
- **R3-P2's prediction lands exactly.** Krit's ATS/HRIS requirement and his dashboards requirement are both MET - EVID citing the named products on L3, in both runs — and the 360 summary line that used to carry them is now NULLED.
- **R3-P4's prediction lands.** All four ^{ss} degree overrides cite tenure and measured impact; **zero** references to the two thresholds v27.7 deleted.
- **Twelve of thirteen verdicts are unchanged from prior.** The exception is Krit: Advance in run A, Advance — Verify in run B.
- **Zero own protected attributes in any of the twenty-six finals.** F11's twelve bias markers grep zero in both runs.

R3-P1 — tool lists admitted

line	records	run A	run B
HubSpot, Zendesk, Metabase, Google Sheets, LINE Official Account, Notion	F1 F2 F3 F4 F7 F8 F11	admitted x7	admitted x7

line	records	run A	run B
SAP WM, Excel, forklift licence (reach truck)	F5 F6	admitted x2	admitted x2
Skills: Workday, Greenhouse, SAP SuccessFactors, ...	Krit L3	admitted	admitted
AdTech Tools: DV360, Appier, GAM, Google Ads, Facebook Ads	Runyarin L29	admitted	admitted
Analytics & Tracking: GA, Adjust, Appsflyer	Runyarin L30	admitted	admitted
any line naming a specific product	Kittiya	<i>unanswerable</i>	<i>unanswerable</i>

Still withheld, as required — generic class words:

line	records	run A	run B
Customer Success \ Onboarding \ Training \ Stakeholder Management \ ...	F2 F4 F6	NULled x3	NULled x3

The boundary the ruling hands to E1 — quote, do not score. Runyarin's EXCEL, POWERPOINT, WORD as bare one-word lines: **admitted, all three, both runs.** E1 v2.7 reads a bare product name on its own line as a named instrument, not as an inventory. It is consistent across runs; whether it is right is the operator's ruling to make.

Side effect worth recording. Under v2.6 the F-battery found the same résumé line NULled on F5 and admitted on F6 (Performed cycle counts and reconciled WMS stock against physical counts). Under v2.7 it is **NULled on both, in both runs** — the pair inconsistency recorded in PROBE-FBAT-REPORT §5.3 is gone.

R3-P2 — what the admitted lists do downstream

Krit — the prediction, met. Both runs, verbatim:

run A:  Experience using an ATS or HRIS such as Workday, Greenhouse, Lever, BambooHR, Manatal, Zoho Recruit, SmartRecruiters, SAP SuccessFactors, or similar: (MET-EVID — Skills line lists Workday, Greenhouse, SAP SuccessFactors, BambooHR, Manatal, Lever)

run A:  Ability to build reports, trackers, or dashboards using Excel, Google Sheets, Looker Studio, Power BI, Tableau, or similar tools: (MET-EVID — Skills line lists Power BI, Tableau, Looker Studio; experience line notes building leadership dashboards that improved hiring efficiency by 40%)

run B:  Experience using an ATS or HRIS such as Workday, Greenhouse, Lever, BambooHR, Manatal, Zoho Recruit, SmartRecruiters, SAP SuccessFactors, or similar: (MET-EVID — Skills list includes Workday, Greenhouse, SAP SuccessFactors, BambooHR, Manatal, Lever)

run B:  Ability to build reports, trackers, or dashboards using Excel, Google Sheets, Looker Studio, Power BI, Tableau, or similar tools: (MET-EVID — Skills list includes Power BI, Tableau, Looker Studio; experience bullet notes leadership dashboards that improved hiring efficiency)

The 360 talent acquisition ... ATS transformation ... HRIS excellence summary line that carried these requirements under v27.6 is **NULled in both runs**. The evidence now comes from the products themselves rather than from a rhetorical line that happened to contain a number.

F1-F11 — no effect, as predicted. Every F-fixture verdict is the one the F-battery recorded under v2.6, where the same line was NULled. A sharper test separates two kinds of product name: Metabase and LINE OA appear both in the skills line and in described work, so citing them proves nothing; HubSpot, Zendesk, Notion, Google Sheets and LINE Official Account appear **only** in the skills line.

product	appears in	citations across the 22 F finals
HubSpot, Zendesk, Notion, Google Sheets, LINE Official Account	skills line only	0
Metabase, LINE OA	skills line <i>and</i> described work	12

Admitting the skills line on a JD that names no tools bought the candidate nothing.

R3-P3 — verdicts

record	prior	R3 A	R3 B	
F1	Advance	Advance	Advance	same
F2	Advance	Advance	Advance	same
F3	Advance (⚡)	Advance	Advance	same
F4	Advance (⚡)	Advance	Advance	same
F5	Do Not Advance	Do Not Advance	Do Not Advance	same
F6	Do Not Advance	Do Not Advance	Do Not Advance	same
F7	Do Not Advance	Do Not Advance	Do Not Advance	same
F8	Advance - Verify	Advance - Verify	Advance - Verify	same
F9	Advance	Advance	Advance	same
F10	Advance	Advance	Advance	same
F11	Advance	Advance	Advance	same
Runyarin	Advance - Verify	Advance - Verify	Advance - Verify	same
Krit	Advance	Advance	Advance - Verify	A≠B

No real applicant is at Do Not Advance. The only **✗** lines in the whole probe are F5's and F6's, the two fixtures built to be declined.

RECORD — Krit is A≠B, on one requirement. *Strong English communication skills for reports and leadership updates:*

run A: **✓** Strong English communication skills for reports and leadership updates: (MET-INF – Resume content describing roles and achievements is written in English, including discussion of leadership dashboards and hiring efficiency)

run B: **⚠** REQ-LANG Strong English communication skills for reports and leadership updates: (verification required – record does not explicitly state English ability; interview should confirm communication level)

Krit's record states no language line at all. The REQ-LANG doctrine in the asset is explicit that a stated language requirement meeting a silent record carries **⚠** REQ-LANG and never blocks — which makes **run B the doctrinally consistent one** and run A's MET-INF, reasoning from the fact that the résumé is written in English, the outlier. Recorded, not scored. It is the difference between an applicant being told to expect a language question at interview and not being told.

R3-P4 — the degree override

All four **⚡** lines on F3 and F4 are quoted in full in §6 of the report. Every one cites tenure **and** measured impact — 46 B2B accounts onboarded, renewal rate improvement from 78% to 91%, time-to-first-report from 21 to 9 days. Grep of both fixtures' finals, both runs, for the two thresholds v27.7 deleted (20 years, 10+ years): **0 hits**. The deleted bullets left no residue, and the reasoning has moved from bare years toward quantified outcomes — which is what the ruling wanted.

Extraction

Both JDs re-extracted under the new E2 hash. **Two distinct JDs, not three:** R3A-FBAT and R3A-S2 share the Happily JD, so it extracts once and S2 cache-hits it inside run A. Four extraction calls in run A, zero in run B.

Both frozen sets: **6 MANDATORY / 0 PREFERRED** (Happily), **8 MANDATORY / 0 PREFERRED** (Bangkok). Zero PREFERRED items, zero Responsibilities items, zero protected tokens in either.

RECORD — the Happily set is not byte-identical to the v27.6 set. Two differences, both punctuation, neither semantic:

- MANDATORY — Bachelor's degree ... {v27.6, U+2019 apostrophe}
- + MANDATORY — Bachelor's degree ... {v27.7, U+0027 apostrophe}

- ... in a customer-facing role with a focus on ... {v27.6}
- + ... in a customer-facing role, with a focus on ... {v27.7, comma added}

Same six requirements, same order, same meaning. Worth knowing because the requirement text is echoed into every evaluation line, so a byte-diff of any R3 evaluation against its F-battery counterpart shows these two differences everywhere.

R3-P6 — invariants

Across all twenty-six finals: L# **0** · Summary **0** · PREFERRED **0** · Second pass? **0** · after the audit identifier **empty** · every E2 pass 1 an evaluation · every E2 pass 2 a full output · header present on all thirteen records in both runs.

F11 bias grep, both runs — นางสาว, Age, 27, 1999, national, Nationality, Single, single, Marital, Buddhist, Religion, Photo: **0 hits**.

City grep across all twenty-six finals: **4 raw hits, all the same string, all JD-sourced**. Every one is the Bangkok JD's own requirement Able to work hybrid in Bangkok 3 days per week — the employer's stated work model, echoed into Krit's requirement preamble and his  NO-CONTRA line. Krit's record states no personal location. Zero hits on any of the twenty-two F-fixture finals. The invariant as written — *own protected attributes in any final: zero* — **holds**; the raw count is reported rather than a filtered one so the master thread sees what the grep saw.

Second pass? still trails every E1 pass-1 output and so reaches E2 inside the NULL list, as PROBE-FBAT-REPORT §9.1 recorded. It never reaches an evaluation.

Calls and tokens

run dir	attempts	failed	total tokens
PROBE-R3A-FBAT	46	0	514,141
PROBE-R3A-S2	4	0	116,340
PROBE-R3A-S3-bkk	6	0	53,292
PROBE-R3B-FBAT	44	0	457,316
PROBE-R3B-S2	4	0	113,365
PROBE-R3B-S3-bkk	4	0	42,834

108 calls against the brief's ≈146. The shortfall is Kittiya's absence (4 calls × 2 runs), the Happily JD extracting once rather than twice in run A (2 calls), and the brief's own arithmetic: a record costs four engine calls — E1 pass 1 and 2, E2 pass 1 and 2 — not five.

One deviation, recorded not corrected

Runyarin was supplied as her intake-normalized text. Her source PDF (4235a383-RunyarinL_CVUpdated2025.pdf) is not in the workspace; runs/20260830-070037/R2/normalized.txt is. run.py numbers a .txt with the same number_lines that intake uses, and the numbering it produced was verified **byte-identical** to the

numbering the original intake produced, before the run. Same deviation and same proof as the Happily JD in PROBE-FBAT-REPORT §0.

What to read for the working

PROBE-R3-REPORT.pdf — 145 pp, bookmarked. §0 execution record · §1 setup and the ADD · §2 extraction and the diff against the prior sets · §3 R3-P1 with every NULL list verbatim · §4 R3-P3 · §5 R3-P2 · §6 R3-P4 · §7 R3-P5 pairs · §8 R3-P6 invariants · §9 A-vs-B diffs · §10 tokens per call · §11 every artifact, both passes, all thirteen records, both runs.

PROBE-R3-EVIDENCE.zip — the six run directories exactly as run.py wrote them, the fixtures, both JDs as fed to the engines, the frozen requirement sets, config.json, manifest.json, the two added assets and their diffs, and FRONTEND-FIX-01-honorific-strip.md untouched.